

Report

HW1-1

(1) Simulate a Function

Function $\sin(5x)/5x$

Model 1

Layer (type)	Output Shape	Param #
Linear-1	$[-1, 1, 1]$	2
ReLU-2	$[-1, 1, 1]$	0
Linear-3	$[-1, 1, 190]$	380
ReLU-4	$[-1, 1, 190]$	0
Linear-5	$[-1, 1, 1]$	191
Total params: 573		
Trainable params: 573		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.00		
Params size (MB): 0.00		
Estimated Total Size (MB): 0.01		


Model 2

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 1]	2
Linear-2	[-1, 1, 5]	10
Linear-3	[-1, 1, 10]	60
Linear-4	[-1, 1, 10]	110
Linear-5	[-1, 1, 10]	110
Linear-6	[-1, 1, 10]	110
Linear-7	[-1, 1, 10]	110
Linear-8	[-1, 1, 5]	55
Linear-9	[-1, 1, 1]	6


Total params: 573
 Trainable params: 573
 Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 0.00
 Params size (MB): 0.00
 Estimated Total Size (MB): 0.00

Loss Chart



Function Curve



Comment:

According to the chart and the graph, we can see that obviously the deep network is better than the shallow network in simulate function task.

(2) Train on Actual Task


Model1

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	1,216
BatchNorm2d-2	[-1, 16, 32, 32]	32
ReLU-3	[-1, 16, 32, 32]	0
MaxPool2d-4	[-1, 16, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	12,832
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 34,634		
Trainable params: 34,634		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.61		
Params size (MB): 0.13		
Estimated Total Size (MB): 0.75		


Model2

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	1,216
BatchNorm2d-2	[-1, 16, 32, 32]	32
ReLU-3	[-1, 16, 32, 32]	0
Conv2d-4	[-1, 16, 32, 32]	6,416
BatchNorm2d-5	[-1, 16, 32, 32]	32
ReLU-6	[-1, 16, 32, 32]	0
MaxPool2d-7	[-1, 16, 16, 16]	0
Conv2d-8	[-1, 8, 16, 16]	3,208
BatchNorm2d-9	[-1, 8, 16, 16]	16
ReLU-10	[-1, 8, 16, 16]	0
Conv2d-11	[-1, 8, 16, 16]	1,608
BatchNorm2d-12	[-1, 8, 16, 16]	16
ReLU-13	[-1, 8, 16, 16]	0
Conv2d-14	[-1, 8, 16, 16]	1,608
BatchNorm2d-15	[-1, 8, 16, 16]	16
ReLU-16	[-1, 8, 16, 16]	0
Linear-17	[-1, 10]	20,490
Total params: 34,658		
Trainable params: 34,658		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.92		
Params size (MB): 0.13		
Estimated Total Size (MB): 1.07		

Loss Chart



ACC Graph



Comment:


According to the result , we can see that the deep cnn performs better than the shallow cnn. From this we can see that building a deeper network can improve the fitting ability of the model.

HW1-2 Optimizer


(1) Visualize the Optimization Process

Collect the weights every 5 epochs, and train 8 times. Reduce the dimension of weights to 2 by PCA

Whole Model




Layer 1



(2) Observe Gradient Norm During Training


Function $\sin 5x/5x$



(3) What Happened When Gradient is Almost Zero

1. To find the weights of the model when the gradient norm is zero, I changed the objective function. The original function is mse loss and I add L1 norm to this loss function. L1 norm can make the weight of network to zero

2. Figure



HW1-3

1. Fit on random labels

Task : Cifar 10


Learning Rate: 0.001

Optimizer ; Adam


Model Summary:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	1,216
BatchNorm2d-2	[-1, 16, 32, 32]	32
ReLU-3	[-1, 16, 32, 32]	0
MaxPool2d-4	[-1, 16, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	12,832
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 34,634		
Trainable params: 34,634		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.61		
Params size (MB): 0.13		
Estimated Total Size (MB): 0.75		

Loss chart



ACC chart



2. Number of parameters v.s. Generalization

10 models

Model1

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 4, 32, 32]	304
BatchNorm2d-2	[-1, 4, 32, 32]	8
ReLU-3	[-1, 4, 32, 32]	0
MaxPool2d-4	[-1, 4, 16, 16]	0
Conv2d-5	[-1, 4, 16, 16]	404
BatchNorm2d-6	[-1, 4, 16, 16]	8
ReLU-7	[-1, 4, 16, 16]	0
MaxPool2d-8	[-1, 4, 8, 8]	0
Linear-9	[-1, 10]	2, 570
Total params: 3,294		
Trainable params: 3,294		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.13		
Params size (MB): 0.01		
Estimated Total Size (MB): 0.15		

Model 2

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 4, 32, 32]	304
BatchNorm2d-2	[-1, 4, 32, 32]	8
ReLU-3	[-1, 4, 32, 32]	0
MaxPool2d-4	[-1, 4, 16, 16]	0
Conv2d-5	[-1, 8, 16, 16]	808
BatchNorm2d-6	[-1, 8, 16, 16]	16
ReLU-7	[-1, 8, 16, 16]	0
MaxPool2d-8	[-1, 8, 8, 8]	0
Linear-9	[-1, 10]	5, 130
Total params: 6,266		
Trainable params: 6,266		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.15		
Params size (MB): 0.02		
Estimated Total Size (MB): 0.19		

Model 3

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 4, 32, 32]	304
BatchNorm2d-2	[-1, 4, 32, 32]	8
ReLU-3	[-1, 4, 32, 32]	0
MaxPool2d-4	[-1, 4, 16, 16]	0
Conv2d-5	[-1, 16, 16, 16]	1,616
BatchNorm2d-6	[-1, 16, 16, 16]	32
ReLU-7	[-1, 16, 16, 16]	0
MaxPool2d-8	[-1, 16, 8, 8]	0
Linear-9	[-1, 10]	10,250
Total params: 12,210		
Trainable params: 12,210		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.20		
Params size (MB): 0.05		
Estimated Total Size (MB): 0.26		

Model 4

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 4, 32, 32]	304
BatchNorm2d-2	[-1, 4, 32, 32]	8
ReLU-3	[-1, 4, 32, 32]	0
MaxPool2d-4	[-1, 4, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	3,232
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 24,098		
Trainable params: 24,098		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.30		
Params size (MB): 0.09		
Estimated Total Size (MB): 0.41		

Model 5

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 32, 32]	608
BatchNorm2d-2	[-1, 8, 32, 32]	16
ReLU-3	[-1, 8, 32, 32]	0
MaxPool2d-4	[-1, 8, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	6,432
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 27,610		
Trainable params: 27,610		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.41		
Params size (MB): 0.11		
Estimated Total Size (MB): 0.52		

Model 6

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	1,216
BatchNorm2d-2	[-1, 16, 32, 32]	32
ReLU-3	[-1, 16, 32, 32]	0
MaxPool2d-4	[-1, 16, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	12,832
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 34,634		
Trainable params: 34,634		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.61		
Params size (MB): 0.13		
Estimated Total Size (MB): 0.75		

Model 7

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	2,432
BatchNorm2d-2	[-1, 32, 32, 32]	64
ReLU-3	[-1, 32, 32, 32]	0
MaxPool2d-4	[-1, 32, 16, 16]	0
Conv2d-5	[-1, 32, 16, 16]	25,632
BatchNorm2d-6	[-1, 32, 16, 16]	64
ReLU-7	[-1, 32, 16, 16]	0
MaxPool2d-8	[-1, 32, 8, 8]	0
Linear-9	[-1, 10]	20,490
Total params: 48,682		
Trainable params: 48,682		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 1.02		
Params size (MB): 0.19		
Estimated Total Size (MB): 1.21		

Model 8

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	4,864
BatchNorm2d-2	[-1, 64, 32, 32]	128
ReLU-3	[-1, 64, 32, 32]	0
MaxPool2d-4	[-1, 64, 16, 16]	0
Conv2d-5	[-1, 64, 16, 16]	102,464
BatchNorm2d-6	[-1, 64, 16, 16]	128
ReLU-7	[-1, 64, 16, 16]	0
MaxPool2d-8	[-1, 64, 8, 8]	0
Linear-9	[-1, 10]	40,970
Total params: 148,554		
Trainable params: 148,554		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 2.03		
Params size (MB): 0.57		
Estimated Total Size (MB): 2.61		


Model 9

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 128, 32, 32]	9,728
BatchNorm2d-2	[-1, 128, 32, 32]	256
ReLU-3	[-1, 128, 32, 32]	0
MaxPool2d-4	[-1, 128, 16, 16]	0
Conv2d-5	[-1, 64, 16, 16]	204,864
BatchNorm2d-6	[-1, 64, 16, 16]	128
ReLU-7	[-1, 64, 16, 16]	0
MaxPool2d-8	[-1, 64, 8, 8]	0
Linear-9	[-1, 10]	40,970
Total params: 255,946		
Trainable params: 255,946		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 3.66		
Params size (MB): 0.98		
Estimated Total Size (MB): 4.64		


Model 10

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 128, 32, 32]	9,728
BatchNorm2d-2	[-1, 128, 32, 32]	256
ReLU-3	[-1, 128, 32, 32]	0
MaxPool2d-4	[-1, 128, 16, 16]	0
Conv2d-5	[-1, 128, 16, 16]	409,728
BatchNorm2d-6	[-1, 128, 16, 16]	256
ReLU-7	[-1, 128, 16, 16]	0
MaxPool2d-8	[-1, 128, 8, 8]	0
Linear-9	[-1, 10]	81,930
Total params: 501,898		
Trainable params: 501,898		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 4.06		
Params size (MB): 1.91		
Estimated Total Size (MB): 5.99		

Train loss Test Loss Chart



Train ACC test ACC chart




HW1-3


Subtask 1

Batchsize 32 vs 64

Train and Loss Chart




Train and Test Acc Chart




SubTask 2

5 diffenet models with batchsize(64 ,128,256,512,1024)

Train Test Loss Chart



Train Test ACC Chart



Sensitivity Chart

